Out-of-Order SMIPS
6.375 Microarchitecture Design
6 April 2011

Team: David S. Greenberg and Bhaskar Mookerji

Abstract

For our 6.375 final project, we are implementing a out-of-order processor using Tomasulo's algorithm and a reorder buffer. The following extends our high-level design to microarchitecture interfaces and discusses an ROB entry tagging scheme that eliminates the need for content addressable memory.

1   Tomasulo's Algorithm at a Glance

Before getting into the details of Tomasulo's algorithm, let's try to understand it at a high level. An instruction set architecture (ISA) has a finite number of places in which it can store variables. These places are called registers. Tomasulo's algorithm dynamically computes the dependencies between instructions so that instructions that don't have data dependencies can run in parallel, while only the truly data-dependent instructions are run in sequence. This increases the throughput of the system by exploiting instruction level parallelism (ILP). Tomasulo's algorithm can be seen as a dynamic hardware implementation of Static Single Assignment, a compiler technique that does the same dependency calculations in software.

The core idea of Tomasulo's algorithm is that each ISA register has a corresponding physical register. Since instructions are dispatched in order but executed out of order, we can associate the source operands of the instruction with the physical registers corresponding to the ISA registers used in the instruction. The destination register of the instruction, however, is written to a new physical register, and the destination ISA register is updated to be linked to the new physical register. That way, every register is assigned once and only true dependencies affect execution.

Of course, hardware can't be allocated, and so Tomasulo's algorithm allocates the physical registers from a pool and releases registers back into the pool when there are no more instructions using them as a source operand. The details of this are explained later.

2   Problem

Although an elastic pipelined microprocessor is an improvement over an unpipelined processor, its performance suffers due to pipeline stalls. There are three classes of stalls, two of which are eliminated by Tomasulo's algorithm.

The first class of stall is a Write-After-Read stall. This is where we see an instruction A which reads from register r_x followed by an instruction B which writes to the same register r_x. Normally, B's execution is blocked by A since they both need to access the same register. If A

Cycle	1	2	3	4	5	6	7	8
Multiply	fetch	exec 1	exec 2	exec 3	exec 4			
Move		fetch	stall	stall	stall	exec		
Add			fetch	stall	stall	stall	exec 1	exec 2

Table 1: Snippet 1 executed on a simply pipelined microprocessor

Cycle	1	2	3	4	5	6
Multiply	fetch	exec 1	exec 2	exec 3	exec 4	
Move		fetch	stall	stall	stall	exec
Add			fetch	exec 1	exec 2	

Table 2: Snippet 1 executed on a microprocessor with Tomasulo's algorithm

has several cycles of latency, then B will be stalled until A completes execution. Tomasulo's algorithm eliminates this false dependency by writing the result of B to a separate physical register, allowing B to execute in parallel with A.

The second class of stall is a Write-After-Write stall. This is where we see an instruction C which writes to a register r_y followed by an instruction D which writes to the same register r_y. Normally, since both instructions need to write to the same ISA register, D would be stalled until C completed; however, Tomasulo's algorithm allows them to execute in parallel since the instructions would each write to separate physical registers.

The third class of stall is a Read-After-Write stall, or a true data dependency. This is where instruction E writes to register r_z, and then instruction F reads from register r_z. Tomasulo's algorithm does nothing in this case, because since F needs the result of E, it must be stalled until E completes.

Let's look at an example snippet of code to understand an example situation in which Tomasulo's algorithm would help (Snippet 1). In this simplified assembly language, the format of an instruction is op, src_1, $[src_2]$, dst; $src2$ is optional.

```
mul r1, r2, r3
mov r3, r4
add r1, r2, r3
```

Snippet 1: Instruction sequence which would benefit from Tomasulo's algorithm

Let's assume that multiplies take 4 clock cycles, moves take 1 clock cycle, and adds take 2 clock cycles. Also, let's assume that fetching an instruction takes 1 clock cycle and happens in parallel with execution. Then executing these instructions in sequence takes a total of 8 clock cycles–one cycle for the initial fetch, followed by each instruction executing in sequence (Table 1). Since fetches are pipelined, we only see the fetch latency on the initial instruction. Now, if we are using Tomasulo's algorithm, we'll instead see a total runtime of 6 clock cycles. To understand why, we can look at this table 2.

3 High-Level Design

The system design and instruction pipeline for our implementation is described in Figure 1 and 2, respectively. The design extends our existing SMIPSv2 implementation into three components: in-order instruction fetch and decoding, out-of-order execution units, and a commit unit. The instruction fetch and decode stage enqueues and dispatches instructions to reservation stations matching the appropriate execution unit, reordering as necessary to remove the pipelining hazards described in Section 2. When the reservation station contains the appropriate operands for a given instruction, execution proceeds, and results are broadcast on a common data bus to waiting reservation stations and the branch prediction unit. The commit unit is implemented as a reorder buffer, and performs in-order commits to memory and the register file. Taken together with dispatch and register renaming, the reorder buffer and the reservation stations implement Tomasulo's algorithm.

4 Goals and Testing Strategy

4.1 Functional Correctness

We will use the existing SMIPS implementation and toolchain to provide inputs, gather outputs, log rule executions as traces, and benchmark our implementation. Verifying the correctness of the Tomasulo's algorithm implementation will first require modular refinement of fundamental components, such as the reorder buffer and reservation stations, before they are incorporated into our existing SMIPS implementation. To formally verify functional correctness, we can examine traces for two different areas: data consistency invariants and instruction completion/termination/stalling conditions[1].

To further check for bugs, we will unit test with short programs that emphasize particular issues: sequences of dependent instructions, sequences of independent arithmetic instructions, small loops, atypical branch sequences, and memory operations. Verifying the traces for these smaller test programs will simplify execution debugging for the existing benchmarks.

4.2 FPGA Synthesis

We've removed all CAM from the system, and so we believe that we should be able to synthesize modulo routing constraints.

4.3 Performance

Raw performance is not the goal of this project. With this architecture, we might not reach the IPCs of the three-stage pipeline architecture in Lab 4. To demonstrate the improvement

---

[1]This criteria for functional correctness in Tomasulo's algorithm are described in greater detail in Chapter 6 of *Design and Evaluation of a RISC Processor with a Tomasulo Scheduler* at `http://www.kroening.com/diplom/diplom/`

of out-of-order execution in our processor, we will artificially introduce latency into our execution units using pipelined integer arithmetic routines. Theoretical performance and correctness are the primary concerns. Our design should enable future developers to use a modularized Tomasulo's algorithm to hide the latencies of their additional instructions. Unfortunately, there is way to show a strict benefit to the elastic pipelined SMIPS processor, since all ALU ops and multiplies are single cycle on an FPGA anyway, and any worthwhile multicycle operation is a project in and of itself to implement.

5 Microarchitecture Overview

5.1 New System Types

To keep in-order instruction issue and retiring through the reorder buffer, we're introducing some new datatypes.

1. **Reorder Buffer Token.** This is used to know which instruction (and thus its destination register) some piece of data corresponds to (could be an instruction or a result) in the reorder buffer. At instruction issue, this type is generated by a token request at the reorder buffer. Keeping track of this value eliminates the need for associative lookup.

2. **Reservation Station Entry.** This stores the reorder token it corresponds to, the operation to execute, and either the computed operands or the token that will contain their result.

3. **Reorder Buffer Entry.** For ALU ops, stores the destination register, the Maybe#(value), and the epoch of the instruction. For branch ops, stores the same stuff as the ALU ops and, if it was a mispredict, the correct address.

4. **Common Data Bus Packet.** The CDB has the ROB token, execution value, and the Maybe#(mispredict).

6 Pipeline Stage Interfaces

The pipeline stages, implemented as modules, are contained with a single processor module similar to the audio pipeline module from the earlier labs. The pipeline stages, connected by the appropriate FIFOs are:

1. **Instruction Fetch**

2. **Instruction Decode and Dispatch**

3. **Execute**

4. **Writeback and Commit**.

The pipeline module interfaces are encapsulated as get/put servers and include the appropriate FIFOs. The connecting datatypes are described in greater detail in the next section.

7 Architectural Interfaces

1. **Issue and Dispatch Units.**

   Inputs: Instruction FIFO

   Outputs: Reservation Station Entry FIFO to the various execute units' reservation stations.

   Uses: Requests ROB token and places ROB entry, updates register map with the ROB token.

   Rules: If ROB token available and instruction ready, unpack instruction, fill in the known operands, dispatch to a reservation station. If store instruction, wait for ROB to empty and issue.

2. **Reservation Stations.**

   Inputs: Reservation Station Entry FIFO

   Outputs: Send instruction to execute unit

   Rules: Snoops CDB and updating operands of held entries until they're ready, issuing instructions when the operands are ready.

3. **Common Data Bus.**

   Inputs: Put execution result (put with method Action)

   Outputs: execution result (multiple reader get (fanout)): all listeners must check the bus every cycle and never drop a packet

   Potential impl: Using bypass reg. for storing value, every listener must acknowledge before accepting a new value

4. **Reorder Buffer.**

   Inputs: Listening to CDB, side effects of token requests

   Outputs: none

   Rules: graduate last instruction if it's been computed (includes writeback and removal from circular buffer), purge instructions of wrong epoch, issue a mispredict if a branch had been mispredicted, update the register map to use the architectural register for map entries whose token matches graduated instruction

5. **Branch Predictor (BTB).** Available methods:

      Return what it would predict given its current state

      Speculate the next instruction

      Return that it mispredicted and needs to fix it

      Return the epoch.

6. **Execute unit.** This is the same as in the pipeline unit, but must also pass reorder token so that it can be bypassed correctly.

Figure 1: System architecture using out-of-order execution and pipelined integer arthmetic.

Figure 2: Instruction pipeline. Pipeline stage is labeled in-box, unless superseded by an overbox. A component box indicates component (such as a FIFO) between stages. Not included in this diagram is the reservation stations and load/store buffer in the execution stage.